

Face Recognition Vendor Test Ongoing

Performance of Automated Presentation Attack Detection (PAD) Algorithms Concept, Evaluation Plan, and API

VERSION 1.2

Mei Ngan
Patrick Grother
Kayee Hanaoka
Austin Hom
Joyce Yang
*Information Access Division
Information Technology Laboratory*

September 30, 2022

Revision History

Date	Version	Description
March 23, 2022	0.1	Draft document for public comment
August 19, 2022	1.0	Intended final document
September 21, 2022	1.1	Additional information on image and video frame size, video length, frame rate, and compression ratio added in Section 5.3. Hardware specifications for the timing test added in Section 1.5.
September 30, 2022	1.2	Additional information on about still input media, algorithmic behavior, and metrics has been added in Sections 5.3, 5.4, and 5.6.

Table of Contents

1. PAD	4
1.1. SCOPE	4
1.1.1. <i>Physical (Analog) vs. Digital Attacks</i>	4
1.1.2. <i>Hardware vs. Software-based PAD</i>	5
1.1.3. <i>Offline Testing</i>	5
1.2. GENERAL FRVT EVALUATION SPECIFICATIONS	5
1.3. REPORTING	6
1.4. ACCURACY METRICS	6
1.5. TIME LIMITS	7
2. RULES FOR PARTICIPATION	7
2.1. PARTICIPATION AGREEMENT	7
2.2. VALIDATION	7
2.3. NUMBER AND SCHEDULE OF SUBMISSIONS	7
3. DATA STRUCTURES SUPPORTING THE API	7
4. IMPLEMENTATION LIBRARY FILENAME	7
5. API SPECIFICATION	8
5.1. HEADER FILE	8
5.2. NAMESPACE	8
5.3. INPUT MEDIA	8
5.4. ALGORITHMIC BEHAVIOR	9
5.5. DEMOGRAPHIC BIAS	9
5.6. PAD SCORE, CALIBRATION, THRESHOLDS, METRICS CALCULATION	9
5.7. API	10
5.7.1. <i>Interface</i>	10
5.7.2. <i>Initialization</i>	10
5.7.3. <i>Presentation Attack Detection</i>	11
5.7.3.1. Presentation Attack Detection - Impersonation Intent	11
1.	12
5.7.3.2. Presentation Attack Detection - Evasion Intent	12

List of Tables

Table 1 – Software vs. hardware-based PAD	5
Table 2 – Processing time limits in milliseconds, per 1280 x 960 image	6
Table 3 – Initialization	10
Table 4 – Presentation Attack Detection	10

1. PAD

1.1. Scope

A presentation attack (PA), as defined by the ISO/IEC 30107¹ standard on biometric presentation attack detection, is “the presentation of an artefact or of human characteristics to a biometric capture subsystem in a fashion intended to interfere with system policy”. A presentation attack is often launched with the intent of *impersonation* (the user is trying to authenticate as a target identity) or *evasion* (the user is trying to fool the biometric system into not recognizing their true identity). The goals of impersonation include trying to gain positive access privilege as someone else, for example, trying to unlock someone’s cell phone or gain access to a facility. The goals of evasion are typically to conceal one’s true identity to evade recognition from say a watchlist, or to create a separate enrollment under a different name. Biometric systems can potentially be attacked by an unknown number of presentation attack instruments, and the number or type of attack instruments in existence is not well-known. Some examples of known presentation attack instruments include artificial “gummy” fingers², “replay” attacks where the attacker is holding a photo or video of someone’s face to the camera³, and iris photo and contact lens attacks⁴.

Presentation attack of face recognition systems (and the ability to detect it) is an area of high interest given the widespread deployment of face recognition systems, particularly in unmanned/unsupervised and remote enrollment and authentication scenarios.

1.1.1. Physical (Analog) vs. Digital Attacks

Presentation attacks have traditionally been associated with physical, analog artefacts that are presented to the sensor. Examples of this include donning a silicone face mask or holding up a printed photo or tablet display to the camera. More recently, a different form of digital presentation attack has surfaced, often called an injection attack, where the attacker bypasses the camera and a digital image or video is injected into the system by virtual camera software. This is a possibility in those scenarios where the capture cannot be trusted - for example in some mobile phones - because the integrity of the sensor/camera cannot be guaranteed. **FRVT PAD intends to evaluate physical attacks with analog artefacts. Digital injection attacks, while an important aspect of the presentation attack landscape, are out of scope for this evaluation.**




Figure 1 - Analog versus digital attacks

¹ ISO/IEC 30107-1:2016 Information technology — Biometric presentation attack detection — Part 1: Framework

² Tsutomu Matsumoto, Hiroyuki Matsumoto, Koji Yamada, and Satoshi Hoshino "Impact of artificial "gummy" fingers on fingerprint systems", Proc. SPIE 4677, Optical Security and Counterfeit Deterrence Techniques IV, (19 April 2002); <https://doi.org/10.1117/12.462719>

³ <https://mobidev.biz/blog/face-anti-spoofing-prevent-fake-biometric-detection>

⁴ Chaos Computer Club Berlin: Chaos Computer Clubs breaks iris recognition system of the Samsung Galaxy S8 (2017). <https://www.ccc.de/en/updates/2017/iriden>

1.1.2. Hardware vs. Software-based PAD

PAD capabilities generally fall into two categories – software-based and hardware-based PAD. Table 1 summarizes the relevance and applications of software vs. hardware-based PAD. The FRVT PAD test will provide ongoing independent testing of **software-based** facial PAD detection technologies. The evaluation is designed to assess software-based PA detection capability to inform developers and current and prospective end-users. Software-based PAD solutions operate only on the captured imagery.

Software-based PAD has a role in those applications where a commodity or non-biometric camera is used for collection of an image that is then transmitted to a receiving system. For example, if a passport or other ID photo is collected in a retail outlet using a generic portrait camera, the passport issuing agency may check for PA. Likewise in some countries' border control points, the capture subsystem may not include a (hardware) PAD module, instead relying on remote server-side PAD operating solely on the image.

This document establishes an initial concept of operations and an application programming interface (API) for evaluation of software-based algorithms to detect facial presentation attack from still photographs and/or video frames.

Note: Hardware-based PAD solutions are currently out of scope in [FRVT PAD](#). For developers interested in evaluation of hardware-based PAD capabilities, the [DHS Science and Technology Directorate](#) is planning a future technology demonstration to include testing of hardware-based PAD capabilities.

Table 1 – Software vs. hardware-based PAD

	Software-based PAD	Hardware-based PAD
Input	Image/video	Image/video + other non-standardized data or signals sensed by dedicated hardware
Mode of operation	Server-based or cloud-based PAD with non-face-aware capture device; offline PAD in existing/legacy systems	Client or edge-based PAD with dedicated face-aware capture device
Applications	Applications where capture processes and devices are not controlled or cannot be configured to perform PAD	Applications where hardware is controllable/configurable during the capture process to perform PAD

1.1.3. Offline Testing

FRVT PAD will be conducted as an offline evaluation at a NIST facility by applying algorithms to still photos and/or video frames that are sequestered on computers controlled by NIST. Offline evaluations are attractive because they allow uniform, fair, repeatable, and large-scale statistically robust testing. However, they do not capture all aspects of an operational system. Offline tests do not include a live image acquisition component or any interaction with real users. Our approach is adopted to allow evaluation on large datasets and to achieve repeatability where all algorithms are tested against the same evaluation datasets.

1.2. General FRVT Evaluation Specifications

General and common information shared between all Ongoing FRVT tracks are documented in a draft version of the FRVT General Evaluation Specifications document – https://pages.nist.gov/frvt/api/FRVT_common_2.0_draft.pdf. This includes rules for participation, hardware and operating system environment, software requirements, reporting, and common data structures that support the APIs.

1.3. Reporting

For all algorithms that complete the evaluation, NIST will provide performance results back to the participating organizations. NIST may additionally report and share results with partner government agencies and interested parties, and in workshops, conferences, conference papers, presentations and technical reports.

Important: NIST will publish the name of the developer’s organization, the algorithm identifiers, and the performance results with attribution to the developer. Results will be machine generated (i.e. scripted) and will include timing, accuracy and other performance results. These will be provided alongside results from other implementations. Results will be expanded and modified as additional implementations are tested, and as analyses are implemented. Results may be regenerated on-the-fly, usually whenever additional implementations complete testing, or when new analyses are added.

Note: Due to data sensitivities, NIST does not intend on disclosing and describing the presentation attack instruments used in our evaluation. We will report PA detection metrics for each presentation attack instrument (PAI) but without the name or description of the PAI. This is intended to encourage broad PAD effectiveness across unknown PAs, and to discourage tuning to specific attacks. This reflects the operational reality that attackers don’t advertise their methods.

1.4. Accuracy metrics

This test will evaluate algorithmic ability to detect whether an image or a video contains a presentation attack or not. Per established metrics⁵ for assessment of presentation attacks, NIST will compute and report:

- Attack Presentation Classification Error Rate (APCER) – the proportion of presentation attack samples incorrectly classified as bona fide presentation
- Bona Fide Presentation Classification Error Rate (BPCER) – the proportion of bona fide samples incorrectly classified as presentation attack samples
- Attack Presentation Non-Response Rate (APNRR⁶) and Bona Fide Presentation Non-Response Rate (BPNRR) – the proportion of presentation attack and bona fide samples, respectively, that do not generate a response and fail to be processed by the algorithm software, whether it’s elective refusal to process the imagery or an involuntary error. Failure to process events will be logged when the algorithm software returns a non-successful return code from the PAD function, indicating that something went wrong while processing the imagery for PAD.

We intend on reporting the above quantities for various presentation attack species.

We intend on incorporating failure to process events into the calculation of BPCER and APCER. All occurrences of failure to process by an algorithm will be treated as if a presentation attack is detected with the confidence score set to +1.

We will also publish error tradeoff plots (BPCER vs. APCER, parametric on threshold) and other analyses as appropriate.

⁵ International Organization for Standardization: Information Technology – Biometric presentation attack detection – Part 3: Testing and reporting. [ISO/IEC FDIS 30107-3:2017](#), JTC 1/SC 37, Geneva, Switzerland, 2017

⁶ ISO/IEC 30107-3 includes APNRR “proportion of attack presentations using the same PAI species that cause no response at the PAD subsystem or data capture subsystem” to quantify outcomes where the sensor is not even triggered by the presented PA sample. We use APNRR to quantify that for the PAD-subsystem, which here is the algorithm under test.

1.5. Time limits

The elemental functions of the implementations shall execute under the time constraints of Table 2. These time limits apply to the function call invocations defined in Section 5. Assuming the times are random variables, NIST cannot regulate the maximum value, so the time limits are median values. This means that the median of all operations should take less than the identified duration. NIST will publish duration statistics.

The time limits apply per image.

Table 2 – Processing time limits in milliseconds, per 1280 x 960 image

Function	
detectImpersonationPA ()	5000 (1 core)
detectEvasionPA ()	

All timing tests will be measured on Intel(R) Xeon(R) Gold 6248 CPU @ 2.50GHz processors. FRVT tests will not support the use of graphics processing units (GPUs).

2. Rules for participation

2.1. Participation agreement

A participant must properly follow, complete, and submit the [FRVT Participation Agreement](#). This must be done once, either prior or in conjunction with the very first algorithm submission. It is not necessary to do this for each submitted implementation thereafter. **Note:** Organizations that have already submitted a participation agreement for FRVT Ongoing 1:1 do not need to send in a new participation agreement unless the organization updates their cryptographic signing key.

2.2. Validation

All participants must run their software through the provided FRVT PAD validation package prior to submission. The validation package will be made available at <https://github.com/usnistgov/frvt>. The purpose of validation is to ensure consistent algorithm output between the participant's execution and NIST's execution. Our validation set is not intended to provide training or test data.

2.3. Number and Schedule of Submissions

NIST will run an initial evaluation of software PAD implementations starting on November 1st, 2022 and will accept algorithm submissions between November 1st - 22nd, 2022. Developers may submit up to two PAD implementations to NIST by November 22nd, 2022. NIST will report results some weeks later. At that point, NIST will determine and announce whether the FRVT PAD evaluation will proceed on an ongoing basis.

The reason for conducting an initial phase of evaluation on a fixed timetable is because we do not know the capability of solutions on our various PA datasets.

3. Data structures supporting the API

The data structures supporting this API are documented in the FRVT - General Evaluation Specifications document available at – https://pages.nist.gov/frvt/api/FRVT_common_2.0_draft.pdf with corresponding header file named *frvt_structs.h* published at https://github.com/usnistgov/frvt/blob/pad/common/src/include/frvt_structs.h.

4. Implementation Library Filename

The core library shall be named as `libfrvt_pad_<provider>_<sequence>.so`, with

- provider: single word, non-infringing name of the main provider. Example: acme
- sequence: a three digit decimal identifier to start at 000 and incremented by 1 every time a library is sent to NIST. Example: 007

Example core library names: *libfrvt_pad_acme_000.so*, *libfrvt_pad_mycompany_006.so*.

Important: Public results will be attributed with the provider name and the 3-digit sequence number in the submitted library name.

5. API specification

Please note that included with the FRVT PAD validation package (available at <https://github.com/usnistgov/frvt>) will be a “null” implementation of this API. The null implementation has no real functionality but demonstrates mechanically how one could go about implementing this API.

5.1. Header File

The prototypes from this document will be written to a file named **frvt_pad.h** and are currently available to implementers at https://github.com/usnistgov/frvt/blob/pad/pad/src/include/frvt_pad.h.

5.2. Namespace

All supporting data structures will be declared in the `FRVT` namespace. All API interfaces/function calls for this track will be declared in the `FRVT_PAD` namespace.

5.3. Input Media

A single image or a sequence of video frames of a single subject will be provided to the PAD algorithms for detection of a presentation attack. All images will contain a single subject presented to the camera.

Presentation Attack Instruments (PAIs)

Initially, the PAIs used in this test are documented in the public domain. NIST aims to expand the diversity of PAIs used in the evaluation over time to reflect the increase in complexity/diversity of attacks. New datasets added to the test will be run on previously submitted algorithms to ensure fair comparison of results.

Capture Environment

- The camera is in a fixed position.
- Images and videos contain a single, cooperative subject, with neutral expression, approximately centered in the field of view.

Still Images

- Image width and height in pixels are provided as input to the algorithm. Image dimensions range from 640x480 to 5184x3456.
- Only color images be used in the evaluation.
- Still imagery used in the evaluation range from losslessly-compressed PNG files to lossy-compressed JPG files. We intend on running both compressed and uncompressed imagery and reporting performance/effects of both.
-

Video

Video sequences can be variable length across all samples provided. Videos will only contain entirely bona fide attempts or entirely presentation attack attempts. That is, there will not be a scenario where a bona fide and PA exist in the same video sequence. For videos, the frame rate in frames per second will be provided to the algorithms.

- **Video Compression**

- The use of compression is operationally realistic for those applications where PAD operates on received media rather than a live-captured stream direct from the sensor. All videos in this test are compressed with h264, with compression ratios in the range of 30:1 to 50:1.
- **Video capture**
 - Video frame rate is provided as input to the algorithm in frames per second (fps). Frame rates range from 24 fps to 60 fps, with typical frame rate of 24 frames per second.
 - Video capture device information is not provided as input to the algorithm.
 - Video duration ranges from 3 seconds to 30 seconds, with typical duration of 25 seconds.
 - Video frame width and height in pixels are provided to the algorithm. Video frame dimensions range from 1920x1080 to 3840x2160.
 - We intend on breaking out analysis and performance by PAI and image vs. video.

NOTE: The imagery provided in the validation package is used for the sole purpose of validation and stress-testing the software. The imagery is not necessarily representative of actual test data that will be used to evaluate the implementations.

5.4. Algorithmic behavior

- When should algorithms report +1.0 in an impersonation attack?
 - If there is any intent to impersonate/spoof a recognition system.
- When should algorithms report 0 in an impersonation attack?
 - No informative value is contained in the image for the software to make a determination either way.
- When should algorithms report -1.0 in an impersonation attack?
 - The image is entirely consistent with natural human faces in typical portrait imagery collected in a cooperative presentation to a camera.
- When should algorithms report +1.0 in an evasion attack?
 - If there is any intent to undermine a recognition outcome. For example, given a photo, large subject deviations from ISO/ICAO portrait photo compliance such as large pose variations, overly exaggerated facial expressions, obstructions of the view of the face could be considered attempts to undermine recognition outcomes.
- When should algorithms report 0 in an evasion attack?
 - No informative value is contained in the image for the software to make a determination either way.
- When should algorithms report -1.0 in an evasion attack?
 - The image is entirely consistent with natural human faces in typical portrait imagery collected in a cooperative presentation to a camera.

5.5. Demographic Bias

Initially, demographic effects will not be a primary evaluation objective, but if observed, we may report concentration of errors specific to particular groups.

5.6. PAD score, calibration, thresholds, metrics calculation

The `detectionImpersonationPA()` and `detectEvasionPA()` functions should return a score between [-1, 1], where 0 is indicating uncertainty in that the input sample has no informative value. We aim to support applications where continuous thresholds can be set based on empirical trials and to that end, we discourage implementations that quantize PAD scores or concentrate scores at +1.0 or -1.0. The error trade off calculation will use a single threshold swept between -1.0 and +1.0. However, recognizing that human operators may be involved in adjudicating whether a questionable sample is indeed a PAI, we may additionally report the PAD score interval between the highest bona fide

scores and lowest PA scores. The isPA boolean returned from the API functions represents a single decision point on a DET at a particular threshold. Operationally, a PAD system would be configured to operate at a particular threshold.

What would be the size of the evaluation datasets - what levels of precision will you be able to reach for APCER and BPCER?

- We anticipate being able to reliably quantify very low BPCER values (below 0.0001) by using large sets of operational bona fide photos. We plan on reporting APCER @ BPCER for various fixed BPCER values.
- We cannot yet gauge APCER values as we have not completed PA image curation.

5.7. API

5.7.1. Interface

The software under test must implement the interface `Interface` by subclassing this class and implementing each method specified therein.

C++ code fragment	Remarks
<pre> Class PADInterface { public: static std::shared_ptr<Interface> getImplementation(); // Other functions to implement }; </pre>	<p>Factory method to return a managed pointer to the <code>Interface</code> object. This function is implemented by the submitted library and must return a managed pointer to the <code>Interface</code> object.</p>

There is one class (static) method declared in `Interface.getImplementation()` which must also be implemented. This method returns a shared pointer to the object of the interface type, an instantiation of the implementation class. A typical implementation of this method is also shown below as an example.

C++ code fragment	Remarks
<pre> #include <frvt_pad.h> using namespace FRVT_PAD; NullImpl:: NullImpl () { } NullImpl::~ NullImpl () { } std::shared_ptr<Interface> Interface::getImplementation() { return std::make_shared<NullImpl>(); } // Other implemented functions </pre>	

5.7.2. Initialization

Before any presentation attack detection calls are made, the NIST test harness will call the initialization function of Table 3. This function will be called BEFORE any calls to `fork()`⁷ are made. This function must be implemented.

Table 3 – Initialization

⁷ <http://man7.org/linux/man-pages/man2/fork.2.html>

Prototype	ReturnStatus initialize(const std::string &configDir);	
		Input
Description	<p>This function initializes the implementation under test. It will be called by the NIST application before any calls to the presentation attack detection functions of this API. The implementation under test should set all parameters. This function will be called N=1 times by the NIST application, prior to parallelizing M >= 1 calls to any other functions via <code>fork()</code>.</p> <p>This function will be called from a single process/thread.</p>	
Input Parameters	configDir	A read-only directory containing any developer-supplied configuration parameters or run-time data files.
Output Parameters	None	
Return Value	See General Evaluation Specifications document for all valid return code values. This function <u>must</u> be implemented.	

5.7.3. Presentation Attack Detection

PAD implementations are often fielded in applications where the classes of risk are known. For example, in authentication, a primary concern is impersonation. In background checks, the concern is of evasion/concealment. The functions of Table 4 and Table 5 separately evaluate presentation attack detection with the intent of impersonation and evasion, respectively. A single image or a sequence of video frames is provided to the functions for detection of a presentation attack. Both PA imagery and non-PA (bona fide) imagery will be used, which will support measurement of attack presentation classification error rate (APCER) with a bona fide classification error rate (BPCER). Developers must implement one or both of these PAD functions. Multiple instances of the calling application may run simultaneously or sequentially. These may be executed on different computers.

5.7.3.1. Presentation Attack Detection - Impersonation Intent

Table 4 – Presentation Attack Detection – Impersonation Intent

Prototypes	ReturnStatus detectImpersonationPA (
	const Media &suspectedPA,	
	bool &isPA,	
	double &score);	
		Input
		Output
		Output
Description	<p>This function takes a piece of input media containing an image or sequence of video frames and outputs a binary decision on whether the media represents a PA and a "padiness" score on [-1, 1] representing how confident the algorithm is that the piece of media contains a PA. A value of -1 means certainty that the media does not contain a PA, and +1 represents certainty that the media does contain a PA. A value near 0 will indicate uncertainty.</p>	
Input Parameters	suspectedPA	Input media of a single still image, or a sequence of video frames. The media will either contain a PA with the intent of impersonating someone else or a bona fide.
Output Parameters	isPA	True if media contains a PA; False otherwise
	score	<p>A real-valued score on [-1, 1]</p> <p>Developers are cautioned that if software only ever reports a few discrete values, the resulting error-tradeoff characteristic will have steps, such that end-users will not be able to set thresholds that finely target some objective (e.g. BPCER = 0.01). This limitation will not occur if the algorithm emits scores with a continuous distribution.</p>
Return Value	See General Evaluation Specifications document for all valid return code values.	

1.

5.7.3.2. Presentation Attack Detection - Evasion Intent**Table 5 – Presentation Attack Detection – Evasion Intent**

Prototypes	ReturnStatus detectEvasionPA(const Media &suspectedPA, bool &isPA, double &score);		
			Input
			Output
			Output
Description	This function takes a piece of input media containing an image or sequence of video frames and outputs a binary decision on whether the media represents a PA and a "padiness" score on [-1, 1] representing how confident the algorithm is that the piece of media contains a PA. A value of -1 means certainty that the media does not contain a PA, and +1 represents certainty that the media does contain a PA. A value near 0 will indicate uncertainty.		
Input Parameters	suspectedPA	Input media of a single still image, or a sequence of video frames. The media will either contain a PA with the intent of evading or concealing the subject's true identity or a bona fide.	
Output Parameters	isPA	True if media contains a PA; False otherwise	
	score	A real-valued score on [-1, 1] Developers are cautioned that if software only ever reports a few discrete values, the resulting error-tradeoff characteristic will have steps, such that end-users will not be able to set thresholds that finely target some objective (e.g. BPCER = 0.01). This limitation will not occur if the algorithm emits scores with a continuous distribution.	
Return Value	See General Evaluation Specifications document for all valid return code values.		